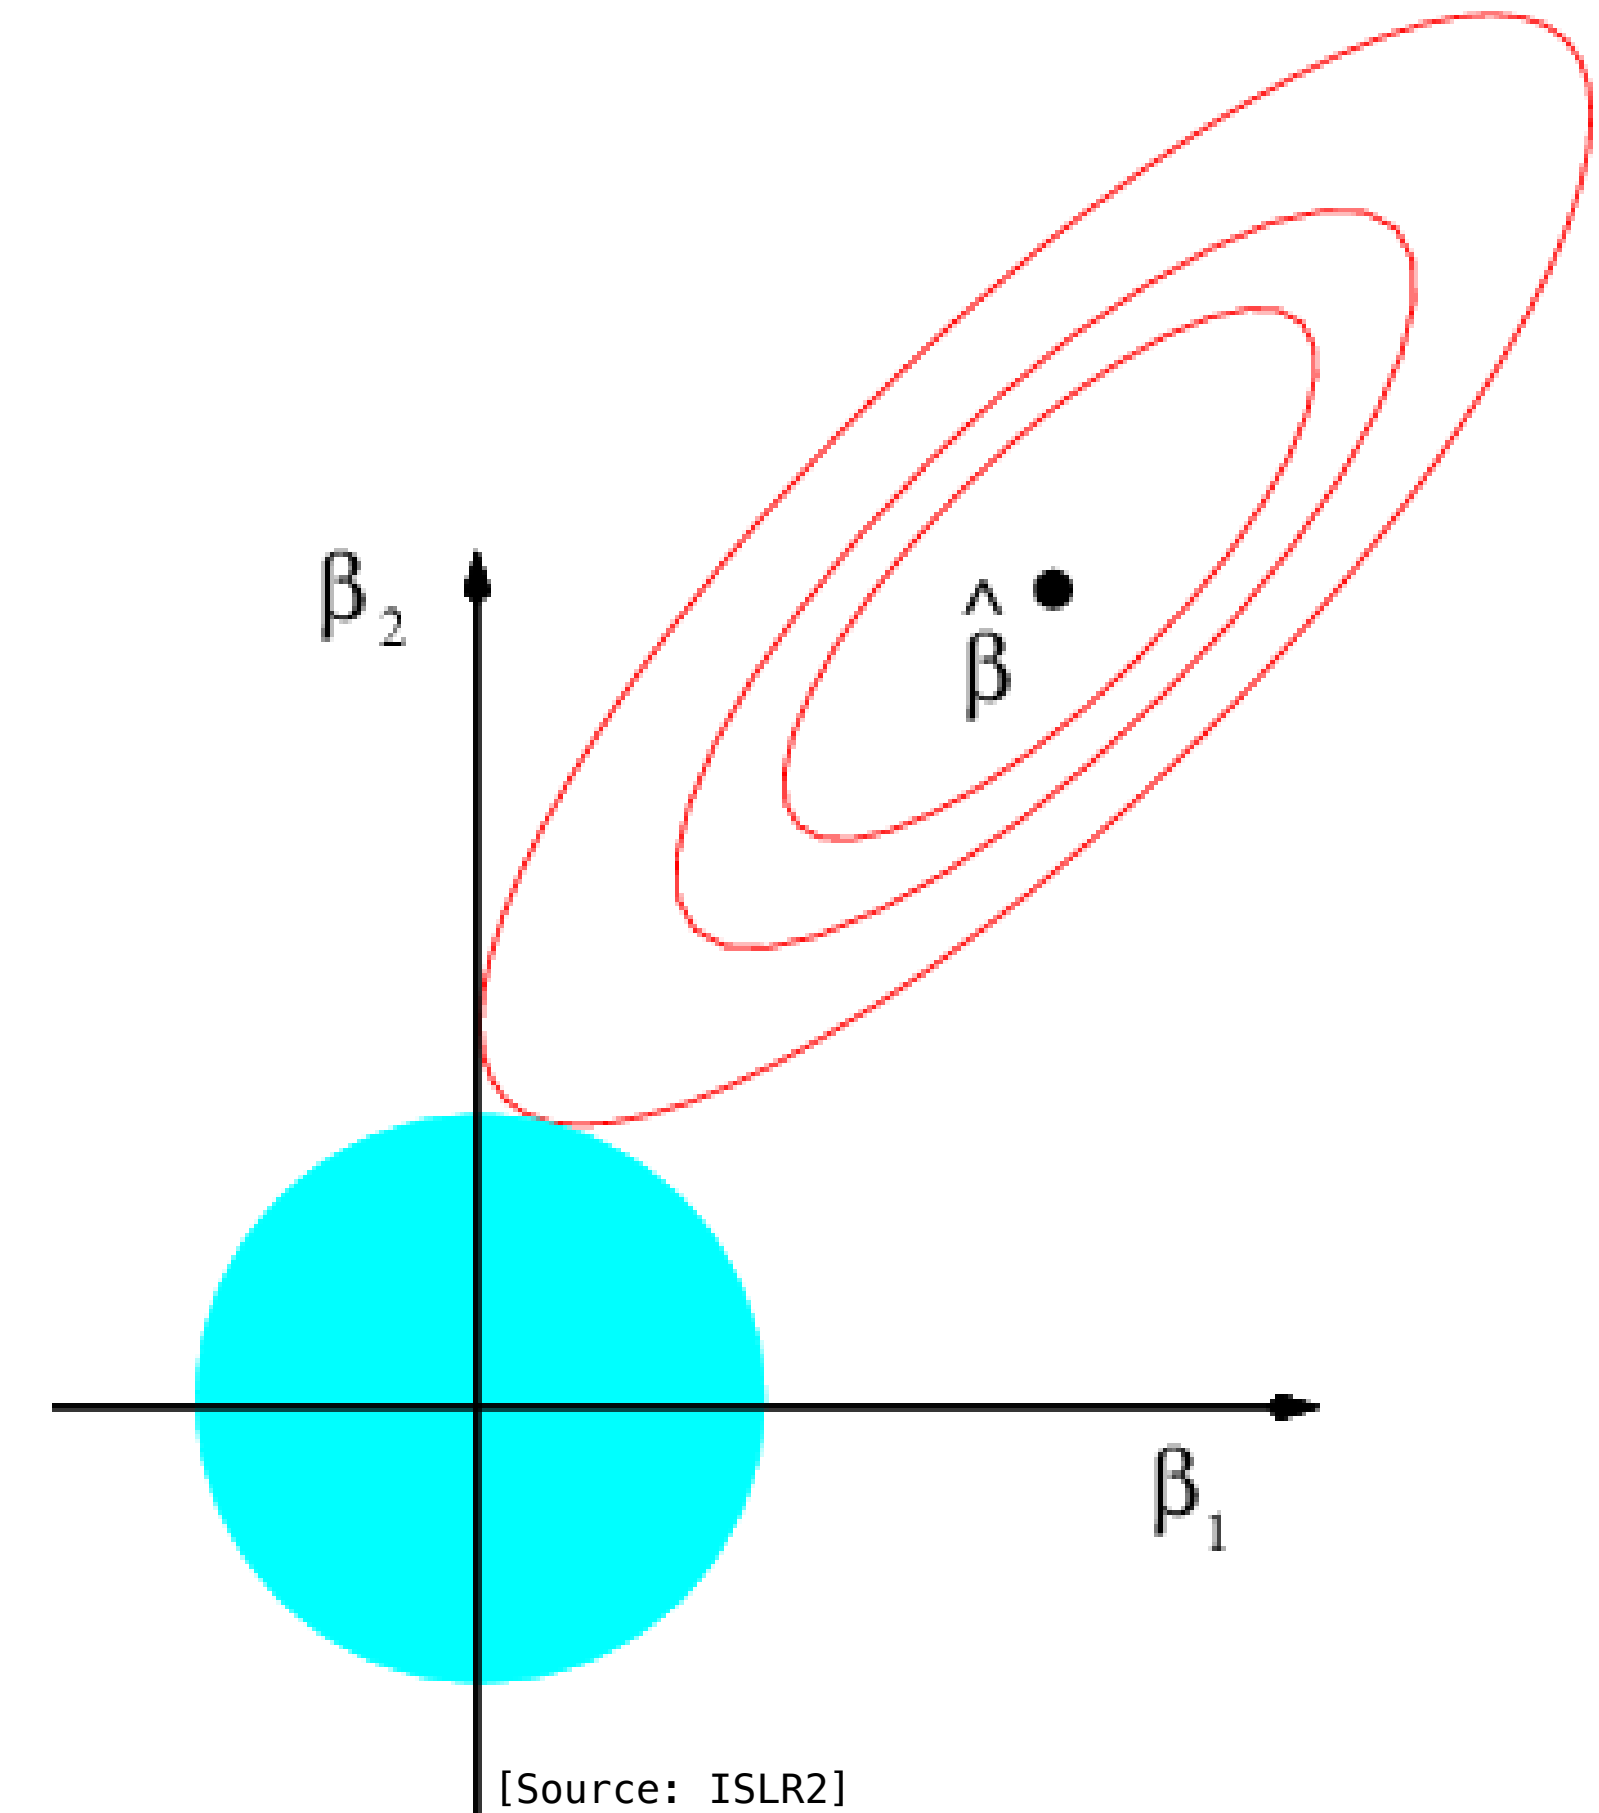# Ridge Regression

ridge uses ℓ_2 penalty

Least Squares produces estimates by minimizing

$$\text{RSS} = \sum_{i=1}^{n} (y_i - \hat{\beta}_0 - \sum_{j-1}^{p} \hat{\beta}_{j_1} x_{ij})^2$$

Ridge regression instead minimizes

$$\underbrace{\sum_{i=1}^{n} (y_i - \hat{\beta}_0 - \sum_{j-1}^{p} \hat{\beta}_{j_1} x_{ij})^2}_{\text{model fit}} + \underbrace{\lambda \sum_{j=1}^{p} \beta_j^2}_{\text{penalty}} = \text{RSS} + \lambda \sum_{j=1}^{p} \beta_j^2$$

[Source: ISLR2]

where $\lambda \geq 0$ is the tuning parameter controlling trade off
between model fit and size of coefficients ($\lambda \to \infty$, $\hat{\beta}_j \to 0$)

Ridge Regression
Regularization Paths

[Source: ISLR2]

$$\ell_2 \text{ norm} = \|\beta\|_2 = \sqrt{\sum_{j=1}^{p} \beta_j^2}$$