Model Selection Criteria

...and compared to cross validation

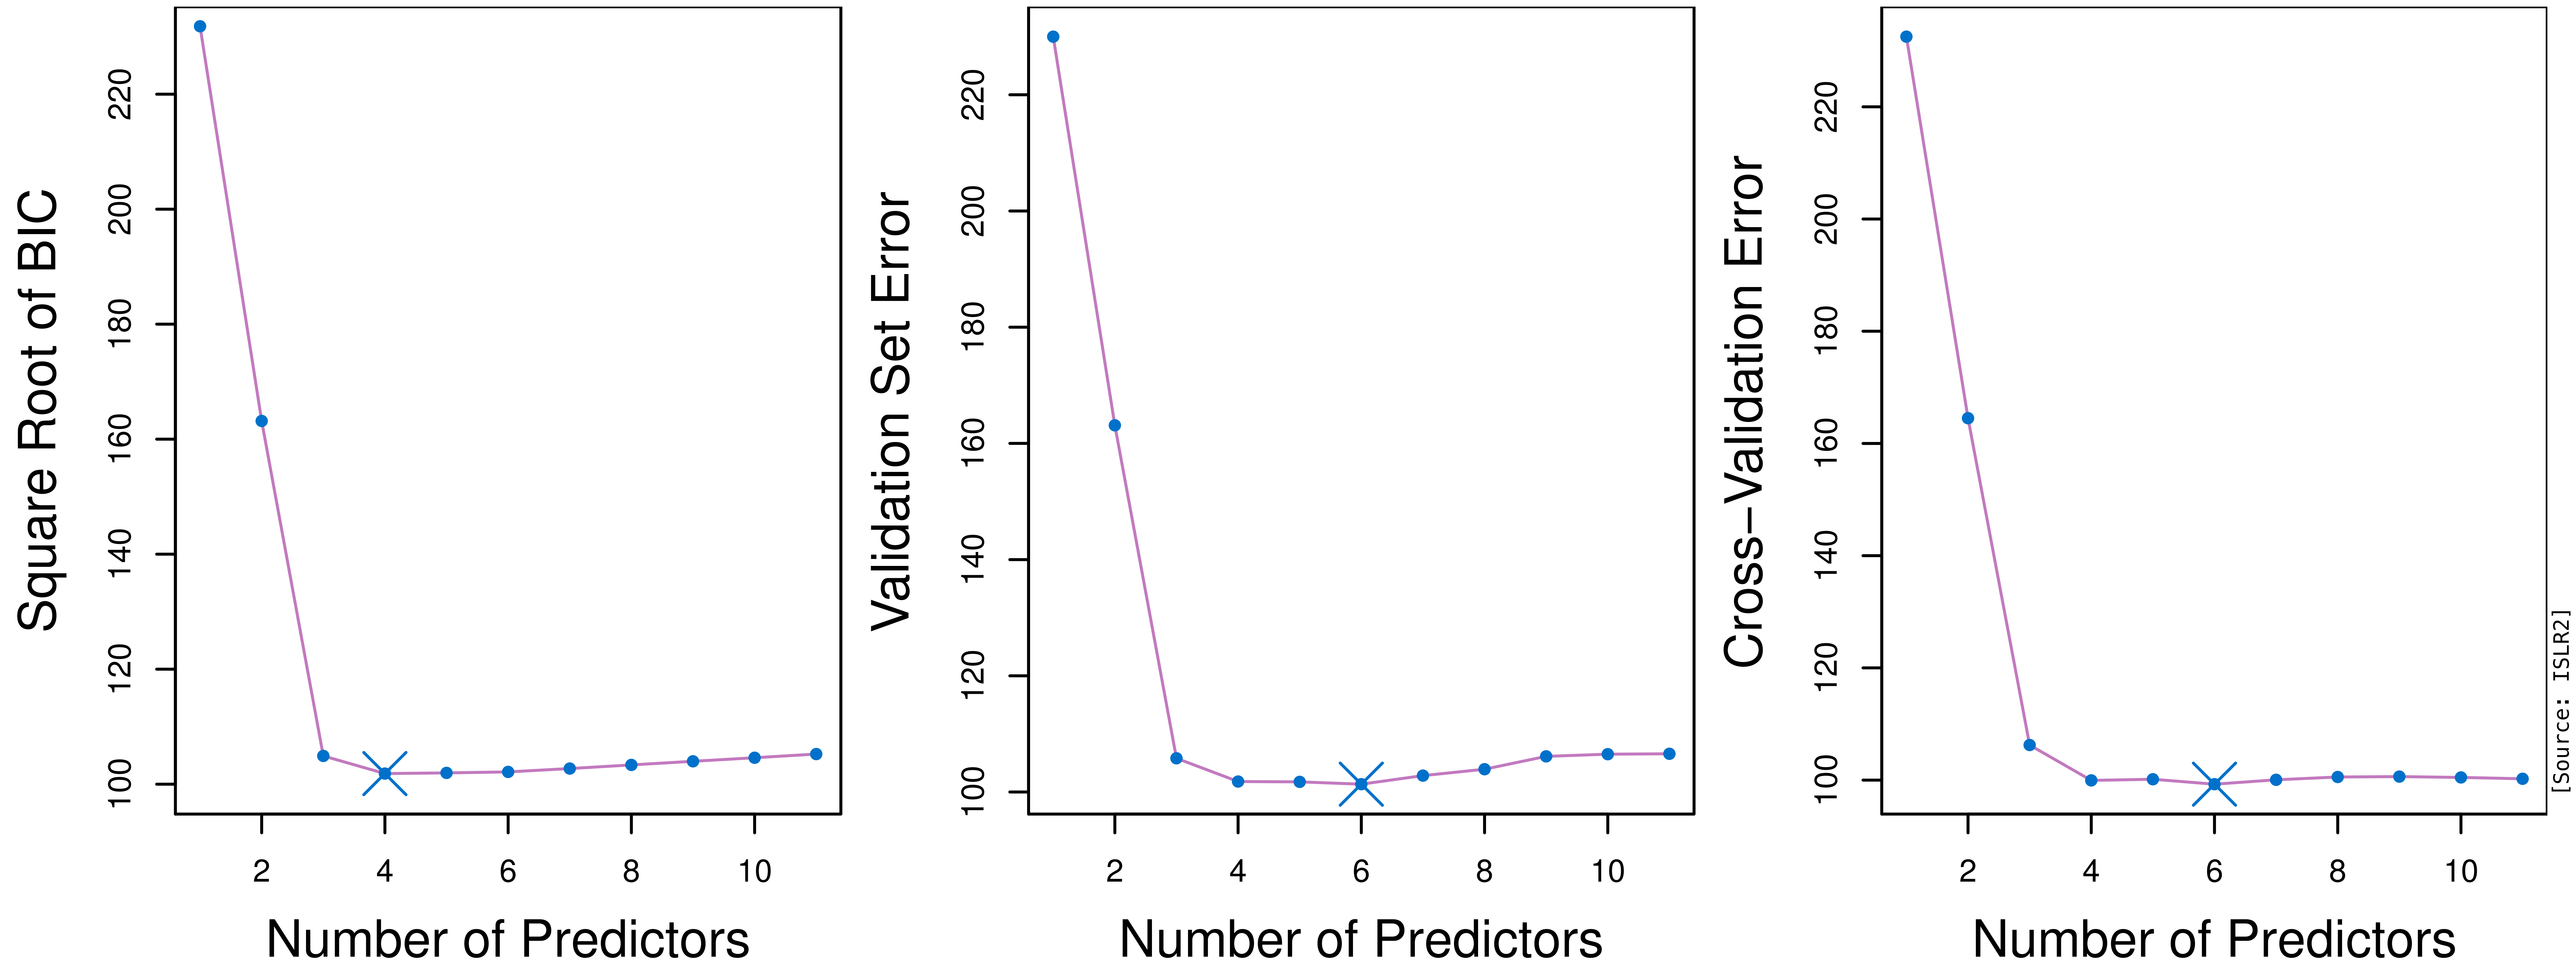

[Source: ISLR2]

Model Search Methods

Best Subset Selection

1. Let M_0 denote null model which contains no predictors. This model simply predicts the response for each observation.
2. For $k = 1,2,\ldots,p$
   - Fit all $\binom{p}{k}$ models that contain exactly p predictors
   - Pick the best among these $\binom{p}{k}$ models and call it M_k.

   Here, *best* is defined as having the smallest RSS or largest R^2
3. Select a single best model from among $M_0, M_1, \ldots, M_p$ using cross validated prediction error, C_p (AIC), BIC, or Adjusted-R^2

requires training 2^p models

Example

$p = 3$

M_0 : intercept only (null)

C_1: X_1 X_2 X_3

lowest training RSS within C_1

$\Longrightarrow M_1$

C_2: X_1, X_2 X_1, X_3 X_2, X_3

lowest training RSS within C_2

$\Longrightarrow M_2$

M_3 : full model with

X_1 X_2 X_3