

- Random forest models have multiple hyper-parameters to tune
 - the number of predictors to randomly select at each split
 - the total number of trees in the ensemble
 - the minimum leaf node size
- Generally tuned through cross validation (\implies data and problem dependent)
- Use **out-of-bag errors** to evaluate model's predictive accuracy
 - cease training once the out-of-bag error stabilizes
 - if sample large enough, estimate is approximately LOO-CV error for bagging
- When the number of predictors is large, but the number of relevant predictors is small, random forests can perform poorly
- Increasing number of trees in ensemble does not increase risk of overfitting
 - but the trees in the ensemble may become more correlated, increase the variance.

Tuning the Random Forest

Tuning the Random Forest

- Random forest models have multiple hyper-parameters to tune
 - the number of predictors to randomly select at each split
 - the total number of trees in the ensemble
 - the minimum leaf node size
- Generally tuned through cross validation (\implies data and problem dependent)
- Use **out-of-bag errors** to evaluate model's predictive accuracy
 - cease training once the out-of-bag error stabilizes
 - if sample large enough, estimate is approximately LOO-CV error for bagging
- When the number of predictors is large, but the number of relevant predictors is small, random forests can perform poorly
- Increasing number of trees in ensemble does not increase risk of overfitting
 - but the trees in the ensemble may become more correlated, increase the variance.

Out-Of-Bag Errors (OOB Errors)

